Data Storage - When reality clashes with theory

Lars Nielsen, *Postdoctoral Researcher*

Aarhus University - Department of Electrical and Computer Engineering

AARHUS UNIVERSITY

Scale IoT

Researchers are often guilty of only celebrating the final success

I want to present the carnage that can happen behind the scene

To create a file system based on a new emerging technology called generalised deduplication

We constantly ran out of disk space

But all analytics tools told us that we had ample space left

So what was wrong?

- Our code
- Something outside our code

- Our code
- Something outside our code

- Our code
  - Full code analysis revealed exactly nothing
  - Follow systems calls, revealed... well, nothing

- Our code
- Something outside our code

- Something outside our code
  - We used EXT4 as an external storage
    - It has something called a directory index

- An index[1] of all files in a folder
- It is (most often) loaded in to memory
- Used to speed up look up operations
  - `ls`, `stat`, etc.

[1] A Directory Index for Ext2, Daniel Phillips, https://www.kernel.org/doc/ols/2002/ols2002-pages-425-438.pdf

- An index of all files in a folder
- It is (most often) loaded in to memory
- Used to speed up look up operations
  - `ls`, `stat`, etc.

So what is the problem?

EXT4 can in theory stored an "unlimited" amount of files

Directory Index has a different limit

EXT4 can in theory stored an "unlimited" amount of files

Directory Index has a different limit

Relax it gets worse the limit varies between Linux distributions and even version.

- Directory Index limits a system that should be limited
- The limit is not well defined

OS	Observed limit
Ubuntu 18.04	12mio
Ubuntu 18.04 server	16mio
Fedora 31	32mio
Fedora 31 Server	32mio
Ubuntu 20.04	20mio
Fedora 33	64mio

- Disable the directory index
  - That is an option
  - It solves the issue
  - But it heavily damages performance

- Disable the directory index
  - That is an option
  - It solves the issue
  - But it heavily damages performance

- So let us hack away around
- We will turn the disadvantage of the directory index into an advantage
- It is only a "solution".

- Our files all have an SHA-1 identifier
  - 20 bytes or converted to hexadecimal string 40 bytes

Let us use that to create a grouping system.

- Our files all have an SHA-1 identifier
  - 20 bytes or converted to hexadecimal string 40 bytes

Let us use that to create a grouping system.

If hash for files shares the 2 first characters, they belong to the same *major group*

and say files that share the first 4 characters belongs to the same *minor group*

- Then, we create a folder for each major group
- In the major group, we create folders for all minor groups
- Then, we place all files belong to that specific minor group in a folder

```
.registry/
└── 00/
    ├── 00/ ..▷ Contains all b where h^b starts with 0000
    ├── 11/ ..▷ Contains all b where h^b starts with 0011
    ├── ...
    └── FF/ ..▷ Contains all b where h^b starts with 00FF
```

What this does is:

- Reduce the probability of hitting the directory index limit
  - all though still present
- But we retain the power of the directory index
- With minimum damage to storage usage -4kB (minimum per folder)
- Zero impact on RAM usage
- Work also for EXT2, EXT3, and ZFS

- Keep a registry in memory of all files stored
- But it increases the RAM consumption of the file system
  - and do you really want your file system to play Google Chrome?

Thank you for your attention